



INTERIM REPORT

RSNA PNEUMONIA DETECTION

Team:

Manjunath S Naragund

GitHub Link: <https://github.com/imanjunathn/Pneumonia-Detection>

Project Summery

In this capstone project, the goal is to build a pneumonia detection system, to locate the position of inflammation in an image. Tissues with sparse material, such as lungs which are full of air, do not absorb the X-rays and appear black in the image. Dense tissues such as bones absorb X-rays and appear white in the image. While we are theoretically detecting “lung capacities”, there are lung capacities that are not pneumonia related. In the data, some of these are labeled “Not Normal No Lung Opacity”. This extra third class indicates that while pneumonia was determined not to be present, there was nonetheless some type of abnormality on the image and oftentimes this finding may mimic the appearance of true pneumonia. Dicom original images: - Medical images are stored in a special format called DICOM files (*.dcm). They contain a combination of header meta data as well as underlying raw image arrays for pixel data.

Problem Statement

In this capstone project, the goal is to build a pneumonia detection system, to locate the position of inflammation in an image. Tissues with sparse material, such as lungs which are full of air, do not absorb the X-rays and appear black in the image. Dense tissues such as bones absorb X-rays and appear white in the image. While we are theoretically detecting “lung capacities”, there are lung capacities that are not pneumonia related. In the data, some of these are labeled “**Not Normal No Lung Opacity**”. This extra third class indicates that while pneumonia was determined not to be present, there was nonetheless some type of abnormality on the image and oftentimes this finding may mimic the appearance of true pneumonia.

Objectives

- ✓ Learn to how to do build an Object Detection Model
- ✓ Use transfer learning to fine-tune a model.
- ✓ Learn to set the optimizer, loss functions, epochs, learning rate, batch size, check pointing, early stopping etc
- ✓ Read different research papers of given domain to obtain the knowledge of advanced models for the given problem.

Methodologies

1. Data set preparation

- Cleaning
- Visualization
- Statistical analysis

2. Image Augmentation

- Image conversion .dcm to .png
- Re sizing as per the model preprocessing requirement

3. Model Building

- Model initialization
- Compiling/Fit the data

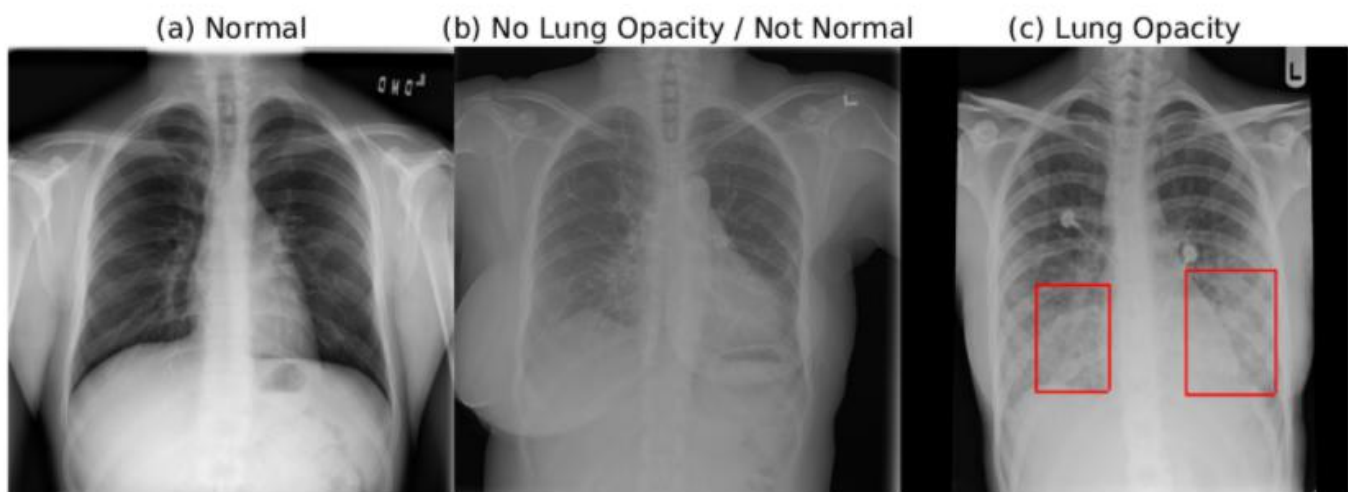
4. Testing accuracy

- Classification report
- Confusion Matrix
- Model performance graph

Dataset / EDA

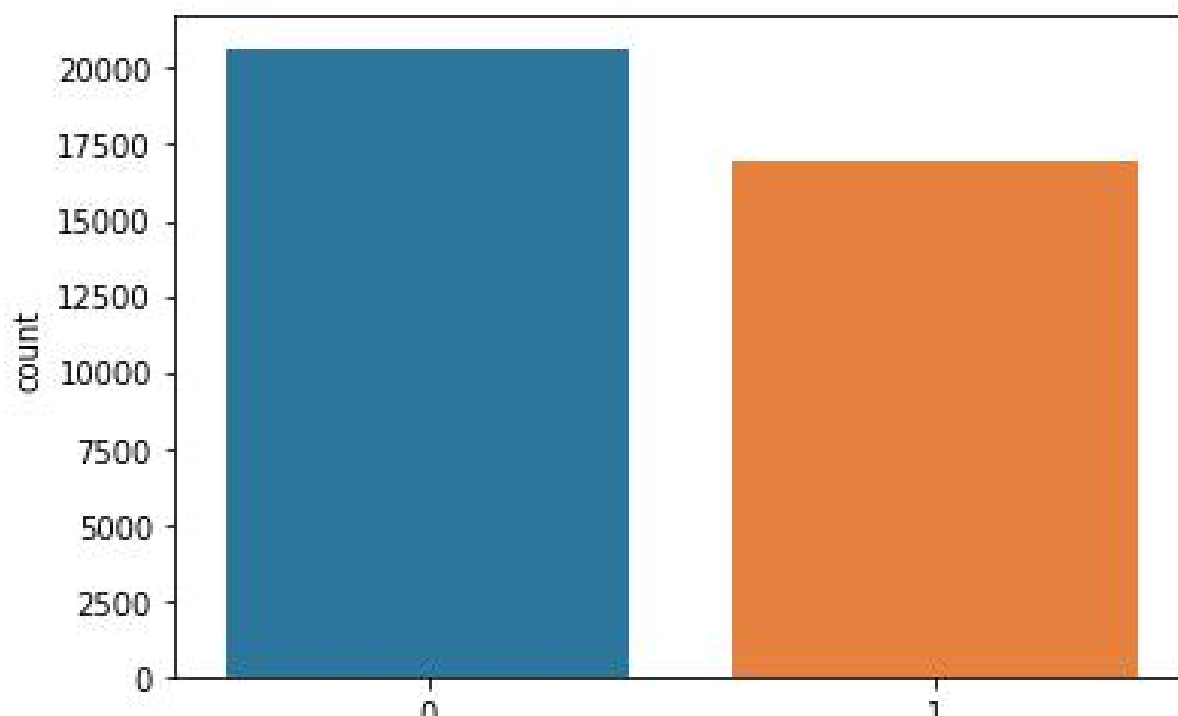
The labeled data-set of the chest X-Ray (CXR) images and patients meta data was publicly provided for the challenge by the US National Institutes of Health Clinical Center. The data-set is available on kaggle platform.

The database comprises frontal-view X-ray images from 26684 unique patients. Each image is labeled with one of three different classes from the associated radiological reports: "Normal", "No Lung Opacity / Not Normal", "Lung Opacity". Fig. 1 shows examples of all three classes CXRs labeled with bounding boxes for unhealthy patients.

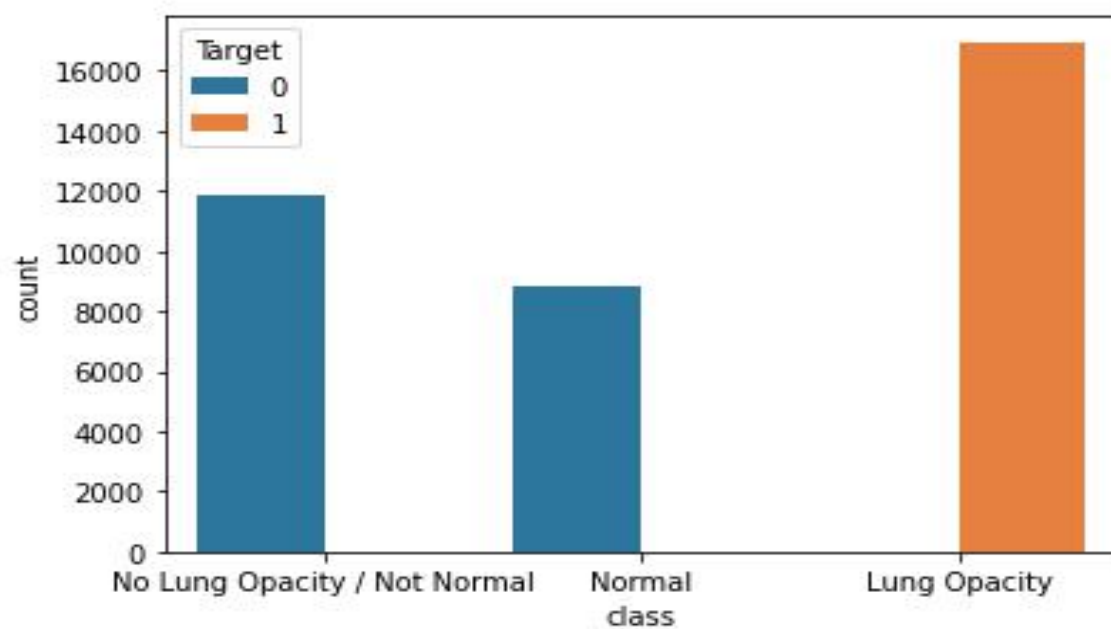


Since we have 3 classes, the distribution among total data is as shown below.

<matplotlib.axes._subplots.AxesSubplot at 0x7fcbbdc4c400>



<matplotlib.axes._subplots.AxesSubplot at 0x7fcbbdbae550>



Checking Distribution of Target

From the CSV data, the target column gives the information of the whether the patient is having pneumonia positive or not by 1 & 0. The distribution of the target is as shown below. Total positive cases are 32% and Negative cases are 68%.

Initial Observations

- Null values are present only with bounding box data I the given CSV.
- All the bounding box null values are associated with target 0
- Each patient ID is associated with a single class or target
- Many of the patient id's are associated with more than one bounding boxes

number_of_patientIDs_per_boxes	
number_of_boxes	
1	23286
2	3266
3	119
4	13



- All positive cases are associated with target 1 only.
- The X-ray images are in .dcm format with a resolution of 1024.
- We have total 30227, in that missing value samples 20672 and 9555.
- We have unique patient id's 26684 and we have duplicate patient id's of 3543.

Sample patient information from the .dcm image is as shown below

```

Dataset.file_meta -----
(0002, 0000) File Meta Information Group Length  UL: 202
(0002, 0001) File Meta Information Version       OB: b'\x00\x01'
(0002, 0002) Media Storage SOP Class UID        UI: Secondary Capture Image Storage
(0002, 0003) Media Storage SOP Instance UID     UI: 1.2.276.0.7230010.3.1.4.8323329.28530.1517874485.775526
(0002, 0010) Transfer Syntax UID                UI: JPEG Baseline (Process 1)
(0002, 0012) Implementation Class UID          UI: 1.2.276.0.7230010.3.0.3.6.0
(0002, 0013) Implementation Version Name       SH: 'OFFIS_DCMTK_360'
-----
(0008, 0005) Specific Character Set             CS: 'ISO_IR 100'
(0008, 0016) SOP Class UID                      UI: Secondary Capture Image Storage
(0008, 0018) SOP Instance UID                   UI: 1.2.276.0.7230010.3.1.4.8323329.28530.1517874485.775526
(0008, 0020) Study Date                        DA: '19010101'
(0008, 0030) Study Time                       TM: '000000.00'
(0008, 0050) Accession Number                  SH: ''
(0008, 0060) Modality                         CS: 'CR'
(0008, 0064) Conversion Type                   CS: 'WSD'
(0008, 0090) Referring Physician's Name        PN: ''
(0008, 103e) Series Description                 LO: 'view: PA'
(0010, 0010) Patient's Name                    PN: '0004cfab-14fd-4e49-80ba-63a80b6bddd6'
(0010, 0020) Patient ID                       LO: '0004cfab-14fd-4e49-80ba-63a80b6bddd6'
(0010, 0030) Patient's Birth Date              DA: ''
(0010, 0040) Patient's Sex                     CS: 'F'
(0010, 1010) Patient's Age                     AS: '51'
(0018, 0015) Body Part Examined                 CS: 'CHEST'
(0018, 5101) View Position                     CS: 'PA'
(0020, 000d) Study Instance UID                 UI: 1.2.276.0.7230010.3.1.2.8323329.28530.1517874485.775525
(0020, 000e) Series Instance UID               UI: 1.2.276.0.7230010.3.1.3.8323329.28530.1517874485.775524
(0020, 0010) Study ID                          SH: ''
(0020, 0011) Series Number                     IS: "1"
(0020, 0013) Instance Number                   IS: "1"
(0020, 0020) Patient Orientation                CS: ''
(0028, 0002) Samples per Pixel                  US: 1
(0028, 0004) Photometric Interpretation         CS: 'MONOCHROME2'
(0028, 0010) Rows                              US: 1024
(0028, 0011) Columns                           US: 1024
(0028, 0030) Pixel Spacing                     DS: [0.14300000000000002, 0.14300000000000002]
(0028, 0100) Bits Allocated                     US: 8
(0028, 0101) Bits Stored                       US: 8
(0028, 0102) High Bit                          US: 7
(0028, 0103) Pixel Representation               US: 0
(0028, 2110) Lossy Image Compression            CS: '01'
(0028, 2114) Lossy Image Compression Method     CS: 'ISO_10918_1'
(7fe0, 0010) Pixel Data                        OB: Array of 142006 elements

```

Steps Followed in Data Preparation

- ✓ Understanding the data with a brief on train/test labels and respective class info
- ✓ Look at the first five rows of both the .csv files(train and test).
- ✓ Identify how are classes and target distributed
- ✓ Check the number of patients with 1, 2, ... bounding boxes
- ✓ Read and extract meta data from dicom files
- ✓ Perform analysis on some of the features from dicom files
- ✓ Check some random images from the training dataset
- ✓ Draw insights from the data at various stages of EDA
- ✓ Visualize some random masks generated

From the dataset, we have classification and regression statement. Whereas the classification part comes with predicting pneumonia positive or negative and the regression part has to predict the area which opacity has found and draw the bounding box.

Image Extraction

Image extraction involves saving image path in input training variable along with making bounding dependency variable & target variable. We handled Null/Nan bounding box variable as zero.

```
def extractImages(foldername,data):
    X_image_train = []
    y_image_train = np.zeros((len(data),4))
    y_train_Target = np.zeros((len(data),1))
    for index,row in data.iterrows():
        name = row[0]
        x1 = int(row[1])
        y1 = int(row[2])
        path = os.path.join(foldername,name)
        path = path+'.png'
        img = cv2.imread(path)
        image_width = img.shape[1]
        image_height = img.shape[0]
        width = int(row[3])
        height = int(row[4])
        target = int(row[5])
        if width != 0 :
            y_image_train[index,0] = x1* image_size/image_width
            y_image_train[index,1] = y1* image_size/image_height
            y_image_train[index,2] = ((width+x1)-x1)* image_size/image_width
            y_image_train[index,3] = ((height+y1)-y1)* image_size/image_width
        else:
            y_image_train[index,0] = 0
            y_image_train[index,1] = 0
            y_image_train[index,2] = 0
            y_image_train[index,3] = 0
        y_train_Target[index] = target
        X_image_train.append(path)

    return (X_image_train,y_image_train,y_train_Target)
```

Image Preprocessing

Pre-processing involves scaling the image to desired size for the selected model. We used PIL package to read, re-size and converting image to RGB(to make 3 channel).

```
def preprocessImage(data):
    processed_data = []
    for i,f in enumerate(data):
        img = Image.open(f)
        img = img.resize((image_size, image_size)) # Resize image
        img = img.convert('RGB')
        processed_data.append(preprocess_input(np.array(img, dtype=np.float32)))
        img.close()
    return processed_data
```

Model Selection

We have prepared a basic model using VGG16 architecture with pre-trained “imagenet” weights. Here we are only training top layers, which is defined by our self. The model has an input shape of image width, height and channel 224, 224, 3 respectively.

The model is fit on 1000 input samples & validation done on 500 samples.

Creating Model

```
1  # Here, I am using VGG16 model
2  # And using input shape (224,224,3) 3 channels.
3  # And using 'imagenet weights'.
4  # This is the basic classification model.
5  def createModel(trainBaseModel=True):
6      inputShape = (image_size,image_size,3)
7      basemodel = VGG16(include_top=False,input_shape=inputShape,weights='imagenet')
8
9      for layer in basemodel.layers:
10         layer.trainable = trainBaseModel
11
12         # basemodel_output = basemodel.get_layer('block5_conv3').output
13         # flat_reg = Flatten()(basemodel_output)
14         # dense = Dense(1024,activation='relu',name='MJ_1_layer_reg')(flat_reg)
15         # drop = Dropout(0.2)(dense)
16         # output_reg = Dense(4,activation='linear',name='output_reg')(drop)
17
18         flat_class = Flatten()(basemodel_output)
19         dense = Dense(1024,activation='relu',name='MJ_1_layer_dense')(flat_class)
20         drop = Dropout(0.2)(dense)
21         output_class = Dense(2,activation='softmax',name='output_class')(drop)
22
23     return Model(inputs=basemodel.input, outputs=[output_class])
```

Note: We are working on classification & we are also trying to implement object detection. Model performance can be increased by using hyper parameter tuning, selecting best optimization algorithm.

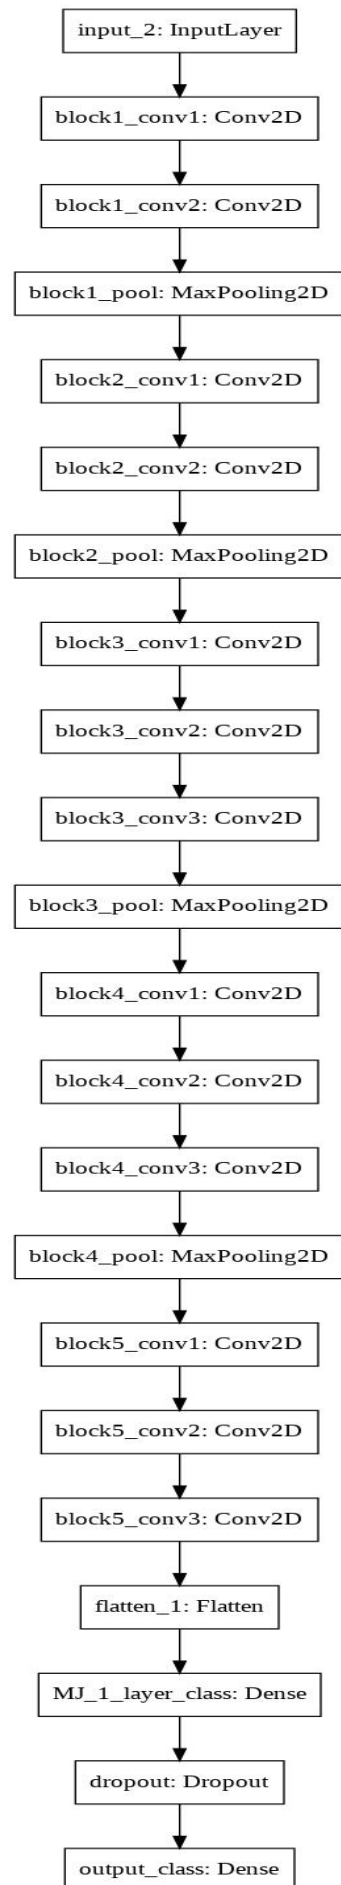
Creating Model

Model: "functional_5"

Layer (type)	Output Shape	Param #	Connected to
=====			
input_7 (InputLayer)	[(None, 224, 224, 3)]	0	
conv2d_14 (Conv2D)	(None, 111, 111, 50)	1400	input_7[0][0]
max_pooling2d_9 (MaxPooling2D)	(None, 55, 55, 50)	0	conv2d_14[0][0]
conv2d_15 (Conv2D)	(None, 53, 53, 100)	45100	max_pooling2d_9[0][0]
max_pooling2d_10 (MaxPooling2D)	(None, 26, 26, 100)	0	conv2d_15[0][0]
conv2d_16 (Conv2D)	(None, 24, 24, 150)	135150	max_pooling2d_10[0][0]
max_pooling2d_11 (MaxPooling2D)	(None, 12, 12, 150)	0	conv2d_16[0][0]
conv2d_17 (Conv2D)	(None, 10, 10, 200)	270200	max_pooling2d_11[0][0]
dropout_11 (Dropout)	(None, 10, 10, 200)	0	conv2d_17[0][0]
flatten_5 (Flatten)	(None, 20000)	0	dropout_11[0][0]
MJ_1_layer_class (Dense)	(None, 1024)	20481024	flatten_5[0][0]
MJ_1_layer_reg (Dense)	(None, 1024)	20481024	flatten_5[0][0]
dropout_13 (Dropout)	(None, 1024)	0	MJ_1_layer_class[0][0]
dropout_12 (Dropout)	(None, 1024)	0	MJ_1_layer_reg[0][0]
output_class (Dense)	(None, 2)	2050	dropout_13[0][0]
output_reg (Dense)	(None, 4)	4100	dropout_12[0][0]
=====			
Total params: 41,420,048			
Trainable params: 41,420,048			
Non-trainable params: 0			

VGG-16 Pneumonia Detection Model

Flow Chart



Model Performance

The model is fit on 1000 input samples & validation done on 500 samples

```
# Use earlystopping
checkpoint = ModelCheckpoint("model-{val_loss:.2f}.h5", monitor="val_loss", verbose=1, save_best_only=True,
                             save_weights_only=True, mode="min", save_freq='epoch')

# Use earlystopping
stop = EarlyStopping(monitor="val_loss", patience=2, mode="min", min_delta=0.01)

model.compile(loss='categorical_crossentropy', optimizer="adam", metrics=['accuracy']) # Regression loss is MSE, Classification: categorical_crossentropy

# Fit the model
history = model.fit(x=X_image_train_temp, y=y_target_train, validation_data=(X_image_test_temp, y_target_test), epochs=10, batch_size=128)
```

Epoch 1/10
8/8 [=====] - 700s 88s/step - loss: 60.2485 - accuracy: 0.7280 - val_loss: 79.3413 - val_accuracy: 0.5960

Epoch 2/10
8/8 [=====] - 695s 87s/step - loss: 8.5367 - accuracy: 0.8910 - val_loss: 9.4288 - val_accuracy: 0.8040

Epoch 3/10
8/8 [=====] - 695s 87s/step - loss: 1.6556 - accuracy: 0.9360 - val_loss: 7.1797 - val_accuracy: 0.7820

Epoch 4/10
8/8 [=====] - 698s 87s/step - loss: 0.3696 - accuracy: 0.9730 - val_loss: 10.5044 - val_accuracy: 0.7100

Epoch 5/10
8/8 [=====] - 698s 87s/step - loss: 0.2221 - accuracy: 0.9830 - val_loss: 6.8278 - val_accuracy: 0.7940

Epoch 6/10
8/8 [=====] - 698s 87s/step - loss: 0.0664 - accuracy: 0.9930 - val_loss: 9.0620 - val_accuracy: 0.7280

Epoch 7/10
8/8 [=====] - 701s 88s/step - loss: 0.0416 - accuracy: 0.9960 - val_loss: 8.9276 - val_accuracy: 0.7300

Epoch 8/10
8/8 [=====] - 697s 87s/step - loss: 5.6399e-04 - accuracy: 1.0000 - val_loss: 7.8051 - val_accuracy: 0.7400

Epoch 9/10
8/8 [=====] - 695s 87s/step - loss: 0.0010 - accuracy: 1.0000 - val_loss: 7.4223 - val_accuracy: 0.7520

Epoch 10/10
8/8 [=====] - 695s 87s/step - loss: 0.0057 - accuracy: 0.9990 - val_loss: 6.6345 - val_accuracy: 0.7660

Model Performance Graph

